

Big Data Engineering

Other Tools and Libraries

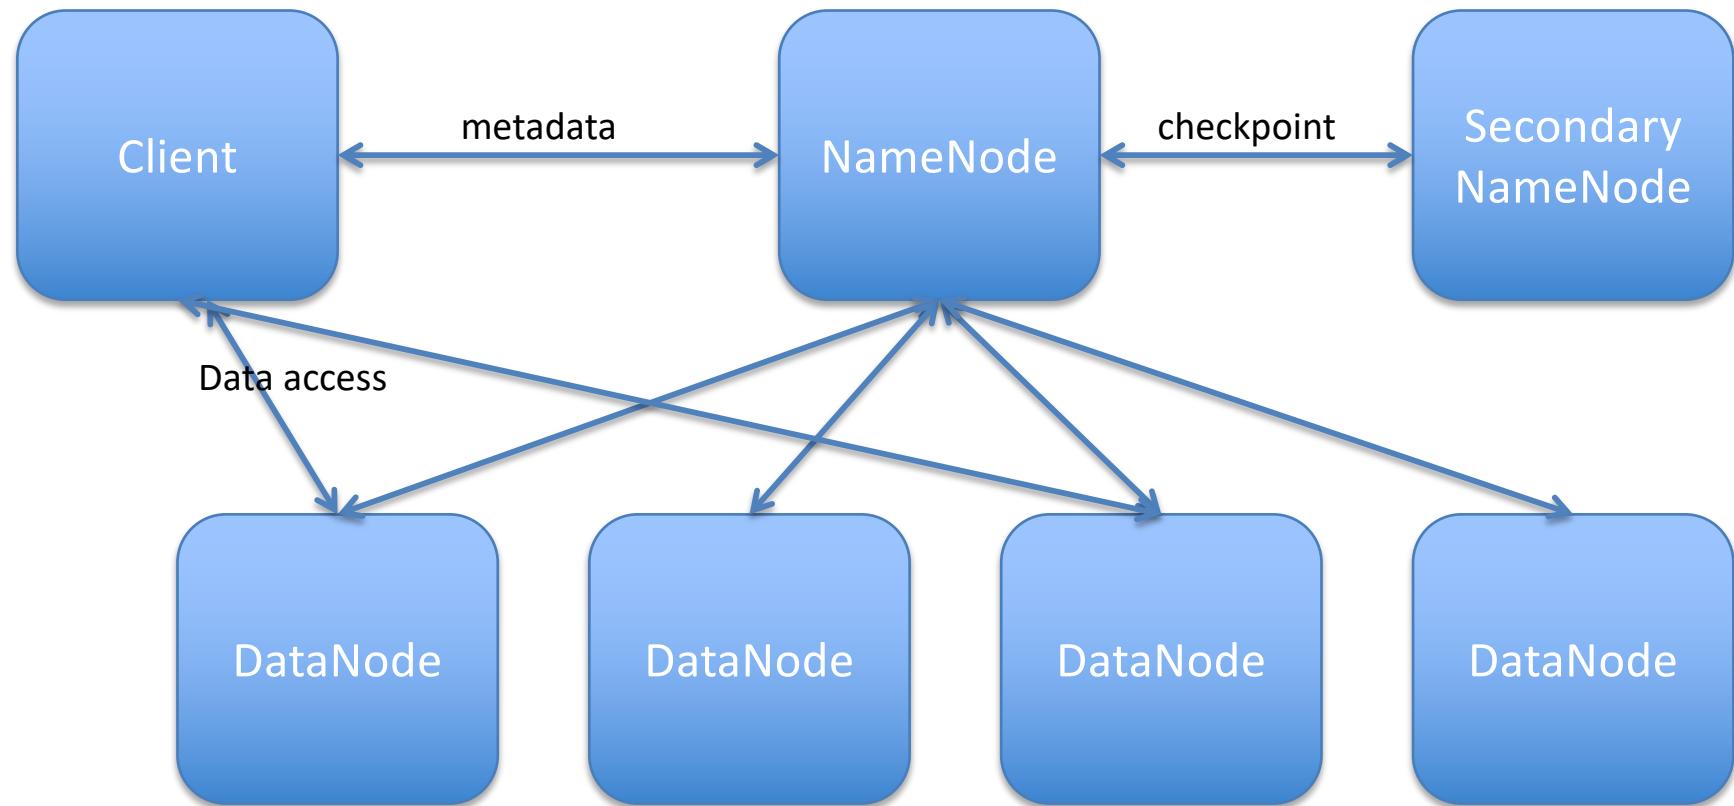
Julie Weeds

March 2019



© Paul Fremantle 2015. This work is licensed under a Creative Commons
Attribution-NonCommercial-ShareAlike 4.0 International License
See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

HDFS in a nutshell



HDFS inspiration

- ## Google File System

- Sanjay Ghemawat, Howard Gobioff, and Shun-Tak Leung. 2003. The Google file system. In Proceedings of the nineteenth ACM symposium on Operating systems principles (SOSP '03). ACM, New York, NY, USA, 29-43.

The Google File System

Sanjay Ghemawat, Howard Gobioff, and Shun-Tak Leung

Google*

ABSTRACT

We have designed and implemented the Google File System, a scalable distributed file system for large distributed data-intensive applications. It provides fault tolerance while running on inexpensive commodity hardware, and it delivers high aggregate performance to a large number of clients.

While sharing many of the same goals as previous distributed file systems, our design has been driven by observations of our application workloads and technological environment, both current and anticipated, that reflect a marked departure from some earlier file system assumptions. This has led us to reexamine traditional choices and explore radically different design points.

The file system has successfully met our storage needs. It is widely deployed within Google as the storage platform

1. INTRODUCTION

We have designed and implemented the Google File System (GFS) to meet the rapidly growing demands of Google's data processing needs. GFS shares many of the same goals as previous distributed file systems such as performance, scalability, reliability, and availability. However, its design has been driven by key observations of our application workloads and technological environment, both current and anticipated, that reflect a marked departure from some earlier file system design assumptions. We have reexamined traditional choices and explored radically different points in the design space.

First, component failures are the norm rather than the exception. The file system consists of hundreds or even thousands of storage machines built from inexpensive com-



© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License
See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

HDFS overview

- Good for streaming access to large files, reliability, scale
- Not good for random access, small files
- Blocks of data 64Mb in size (configurable)
- Each block can be replicated across multiple data nodes for High Availability (HA)



© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License
See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

HDFS Usage

- Spotify has 1600+ nodes, storing 60+ petabytes of data
 - <https://www.usenix.org/system/files/conference/fast17/fast17-niazi.pdf>
- One of Facebook's largest clusters (based on HDFS) holds more than 100 PB of data, processing more than 60,000 Hive queries a day
 - <https://www.facebook.com/notes/facebook-engineering/under-the-hood-scheduling-mapreduce-jobs-more-efficiently-with-corona/>



HopFS

- HopFS is a drop-in replacement for HDFS, based on HDFS v2.0.4.

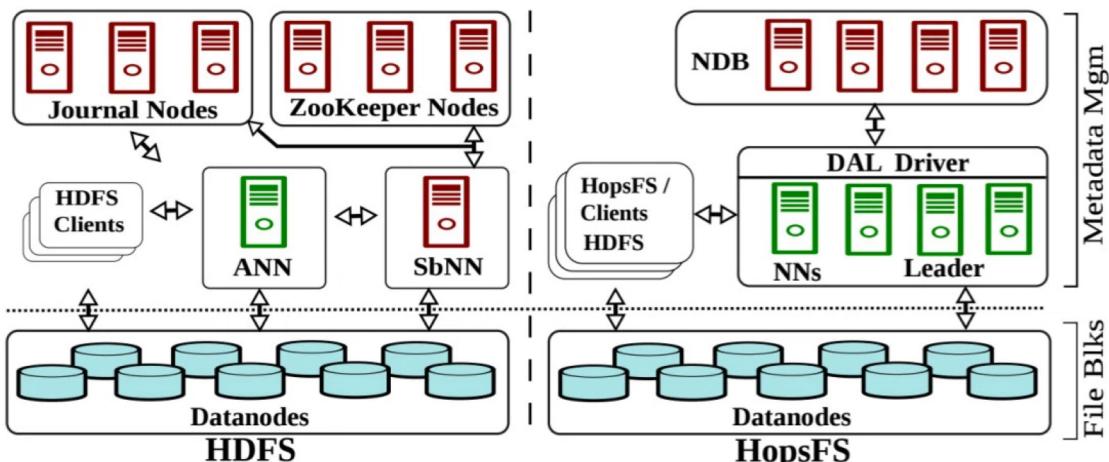


Figure 1: System architecture for HDFS and HopFS. For high availability, HDFS requires an Active NameNode (ANN), at least one Standby NameNode (SbNN), at least three Journal Nodes for quorum-based replication of the write ahead log of metadata changes, and at least three ZooKeeper instances for quorum based coordination. HopFS supports multiple stateless namenodes that access the metadata stored in NDB database nodes.

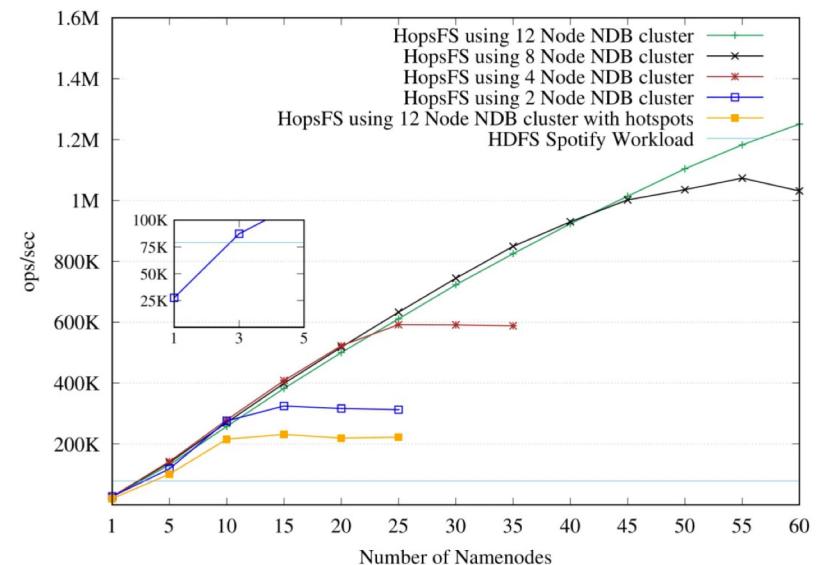
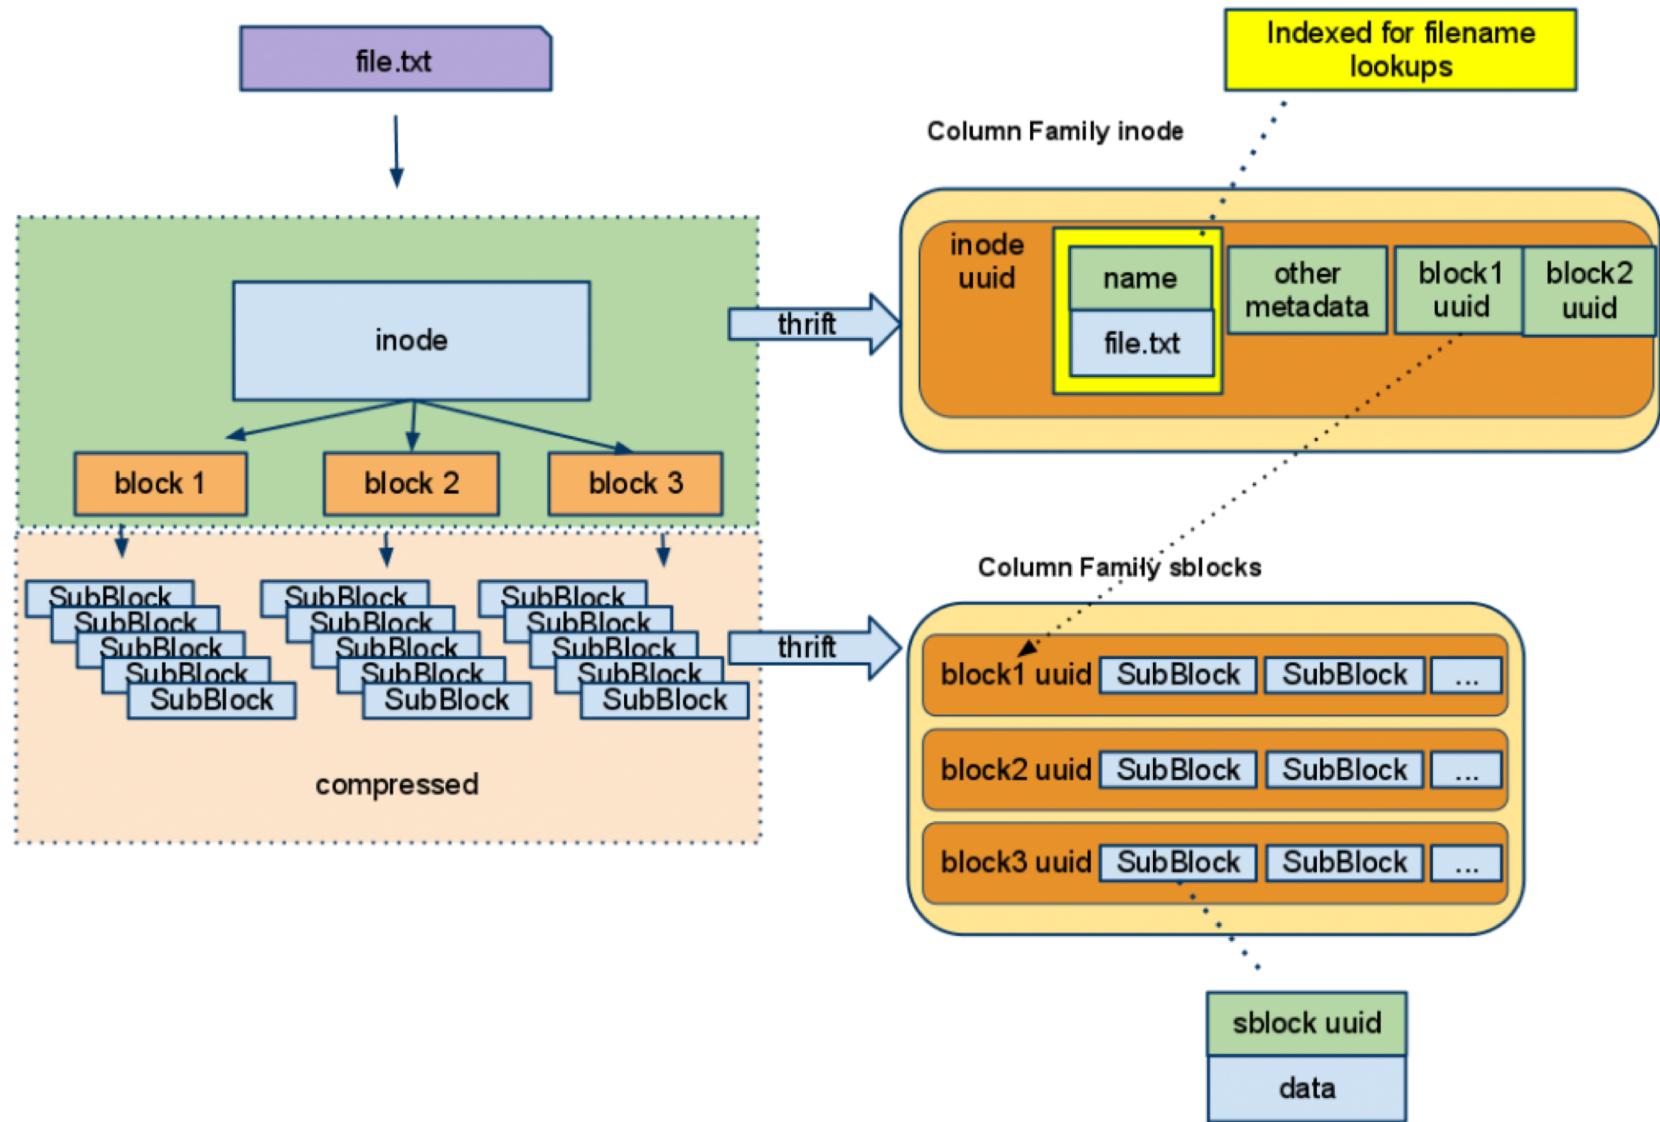


Figure 6: HopsFS and HDFS throughput for Spotify workload.

CassandraFS (not open source)



Amazon S3

- Simple Storage Service
- Unlimited storage of files
 - Up to 5 terabytes each
 - Stored in named “buckets”
 - Accessible via AWS APIs or HTTP
 - Authenticated or Public



© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License
See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

Spark packages

- A wide set of plugins
 - Currently 148 community donated plugins
- Data connectors
 - Cassandra, Couchbase, Mongo, CSV, etc
- Machine Learning, Neural networks
- Streaming
- etc



Using Spark Packages

Automatic download from the web:

bin/spark-shell

--packages com.databricks:spark-csv_2.11:1.2.0



© Paul Fremantle 2015. This work is licensed under a Creative Commons
Attribution-NonCommercial-ShareAlike 4.0 International License
See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

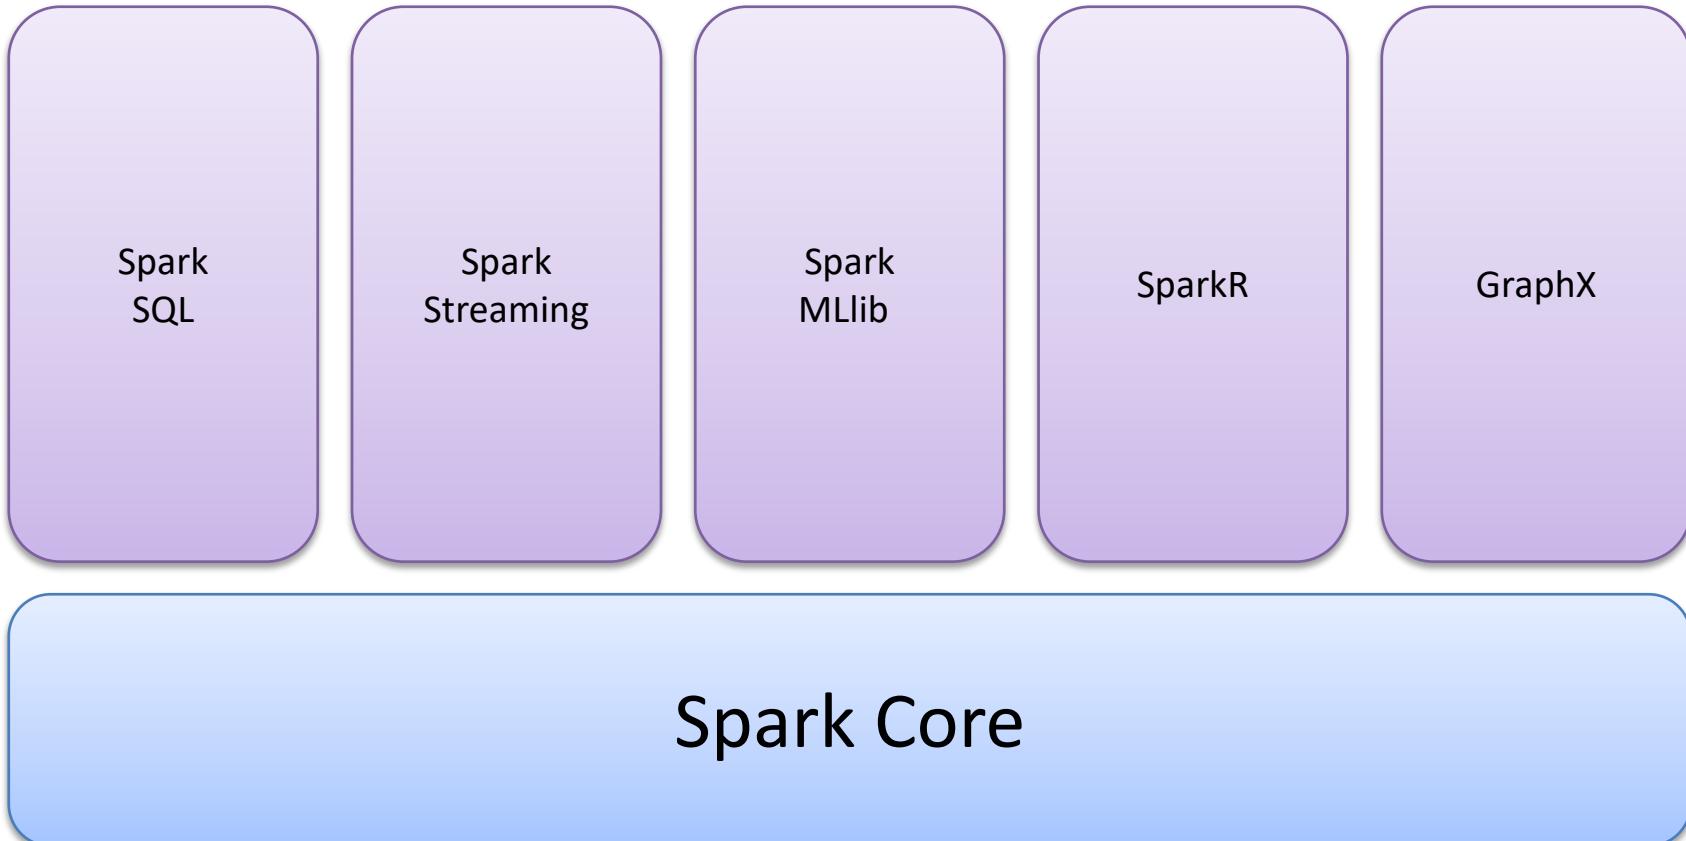
Locality

- Spark understands the locality of data:
 - Already in memory
 - HDFS location
 - Cassandra location



© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License
See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

Spark Extras



Spark Extras

- Spark Streaming
 - Realtime analysis in Spark
- Spark MLlib
 - Like Mahout – Machine learning in Spark
- GraphX
 - Graph processing in Spark
- SparkR
 - R statistical analysis on Spark



© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License
See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

Spark MLlib

- Simple stats and correlation testing
- Classification and regression
- Collaborative Filtering
 - Alternating Least Squares
- Clustering
 - k-means, etc
- Frequent Pattern Mining
- Plus more



MLlib example

```
from pyspark.mllib.fpm import FPGrowth

data = sc.textFile("data/mllib/sample_fpgrowth.txt")

transactions = data.map(lambda line: line.strip().split(' '))

model = FPGrowth.train(transactions, minSupport=0.2,
    numPartitions=10)

result = model.freqItemsets().collect()
for fi in result:
    print(fi)
```

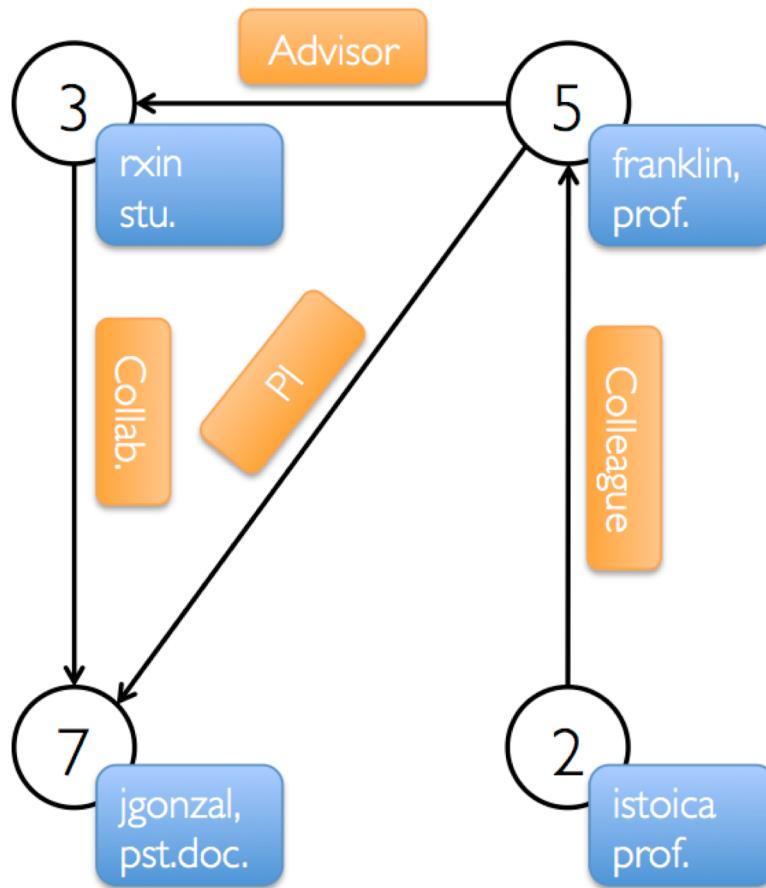




GraphX

GraphX

Property Graph



Vertex Table

Id	Property (V)
3	(rxin, student)
7	(jgonzal, postdoc)
5	(franklin, professor)
2	(istoica, professor)

Edge Table

SrcId	DstId	Property (E)
3	7	Collaborator
5	3	Advisor
2	5	Colleague
5	7	PI



R



- R is an open source system for statistics and graphics
 - Based on the S language from AT&T Bell Labs
- Supports a wide variety of statistical techniques and graphing tools
- An extensible set of packages that provide extra functions via CRAN
 - The Comprehensive R Archive Network



SparkR

- A lightweight approach to use Spark from within R
- Also works with MLlib for machine learning
- Allows complex statistical analysis to be done on a Spark cluster



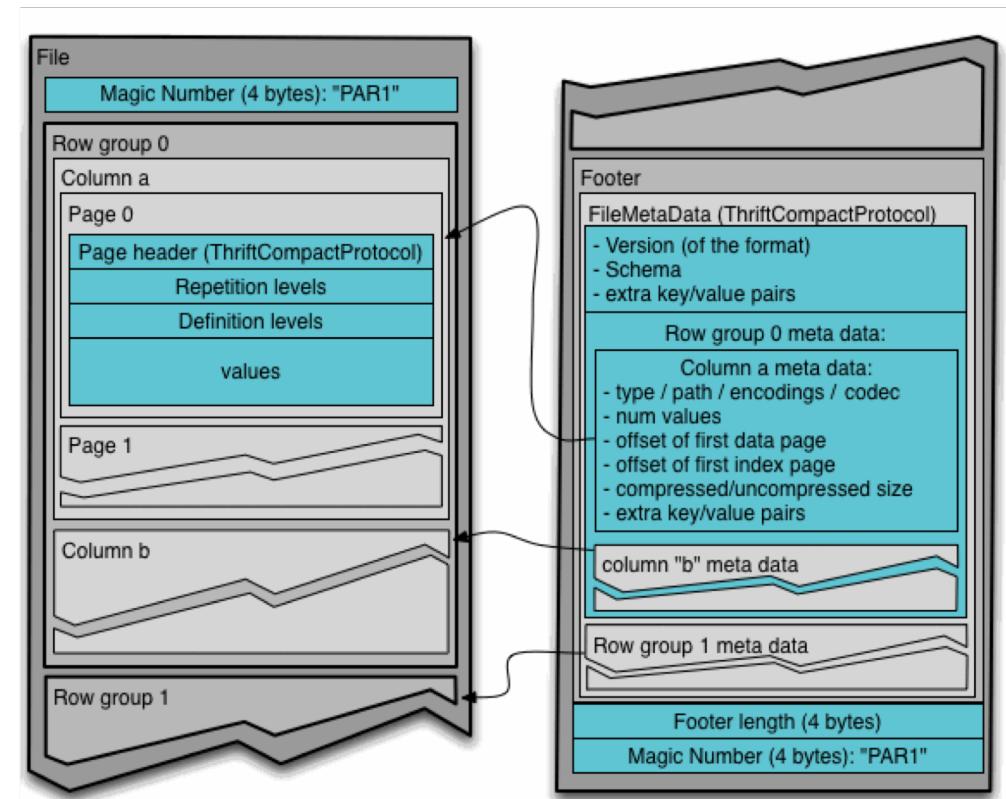
Apache Avro

- A compact data storage and transmission system
 - Uses schemas of data to ensure it can be read by the receiver
 - Supports dynamic typing
- Used by RPC or data collection systems
 - Fast binary protocols
- Also supports storage
 - Hence used by many Big Data apps including Hadoop and Spark



Apache Parquet

- Apache Parquet is a columnar data storage model
 - Works with Hadoop, Spark and many others
 - Efficient storage of data
 - Based on another Google system called Dremel



Cluster management systems for Big Data

- YARN
 - Part of Hadoop but significantly rebuilt since Hadoop 1
- Mesos
 - Popular Apache project
 - Built to be a resource manager for a complete datacenter
 - Supports many workloads (e.g. Docker as well as Spark)

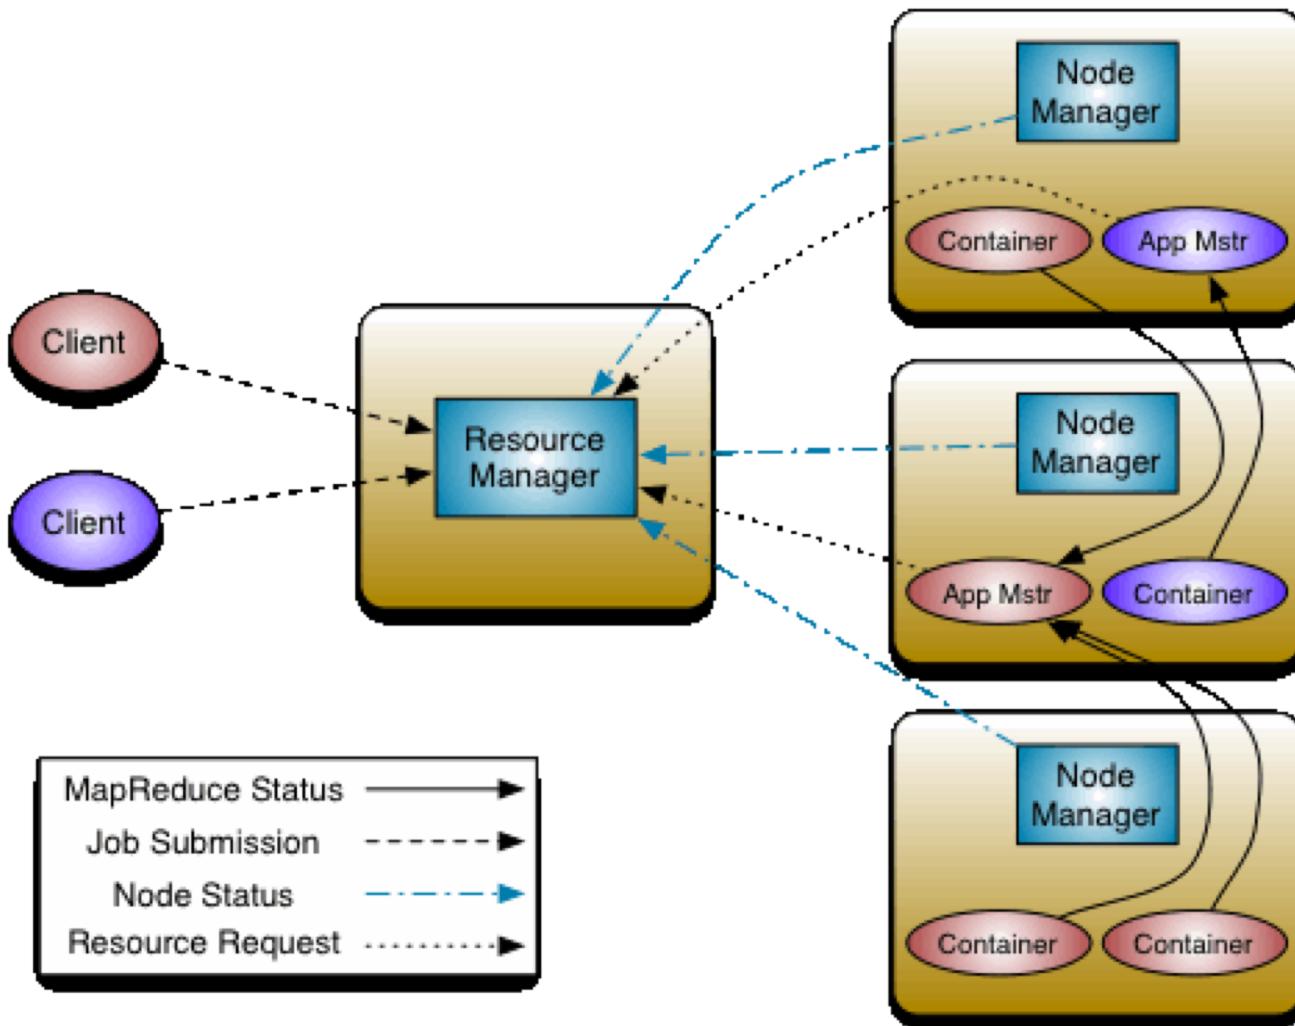


What is YARN?

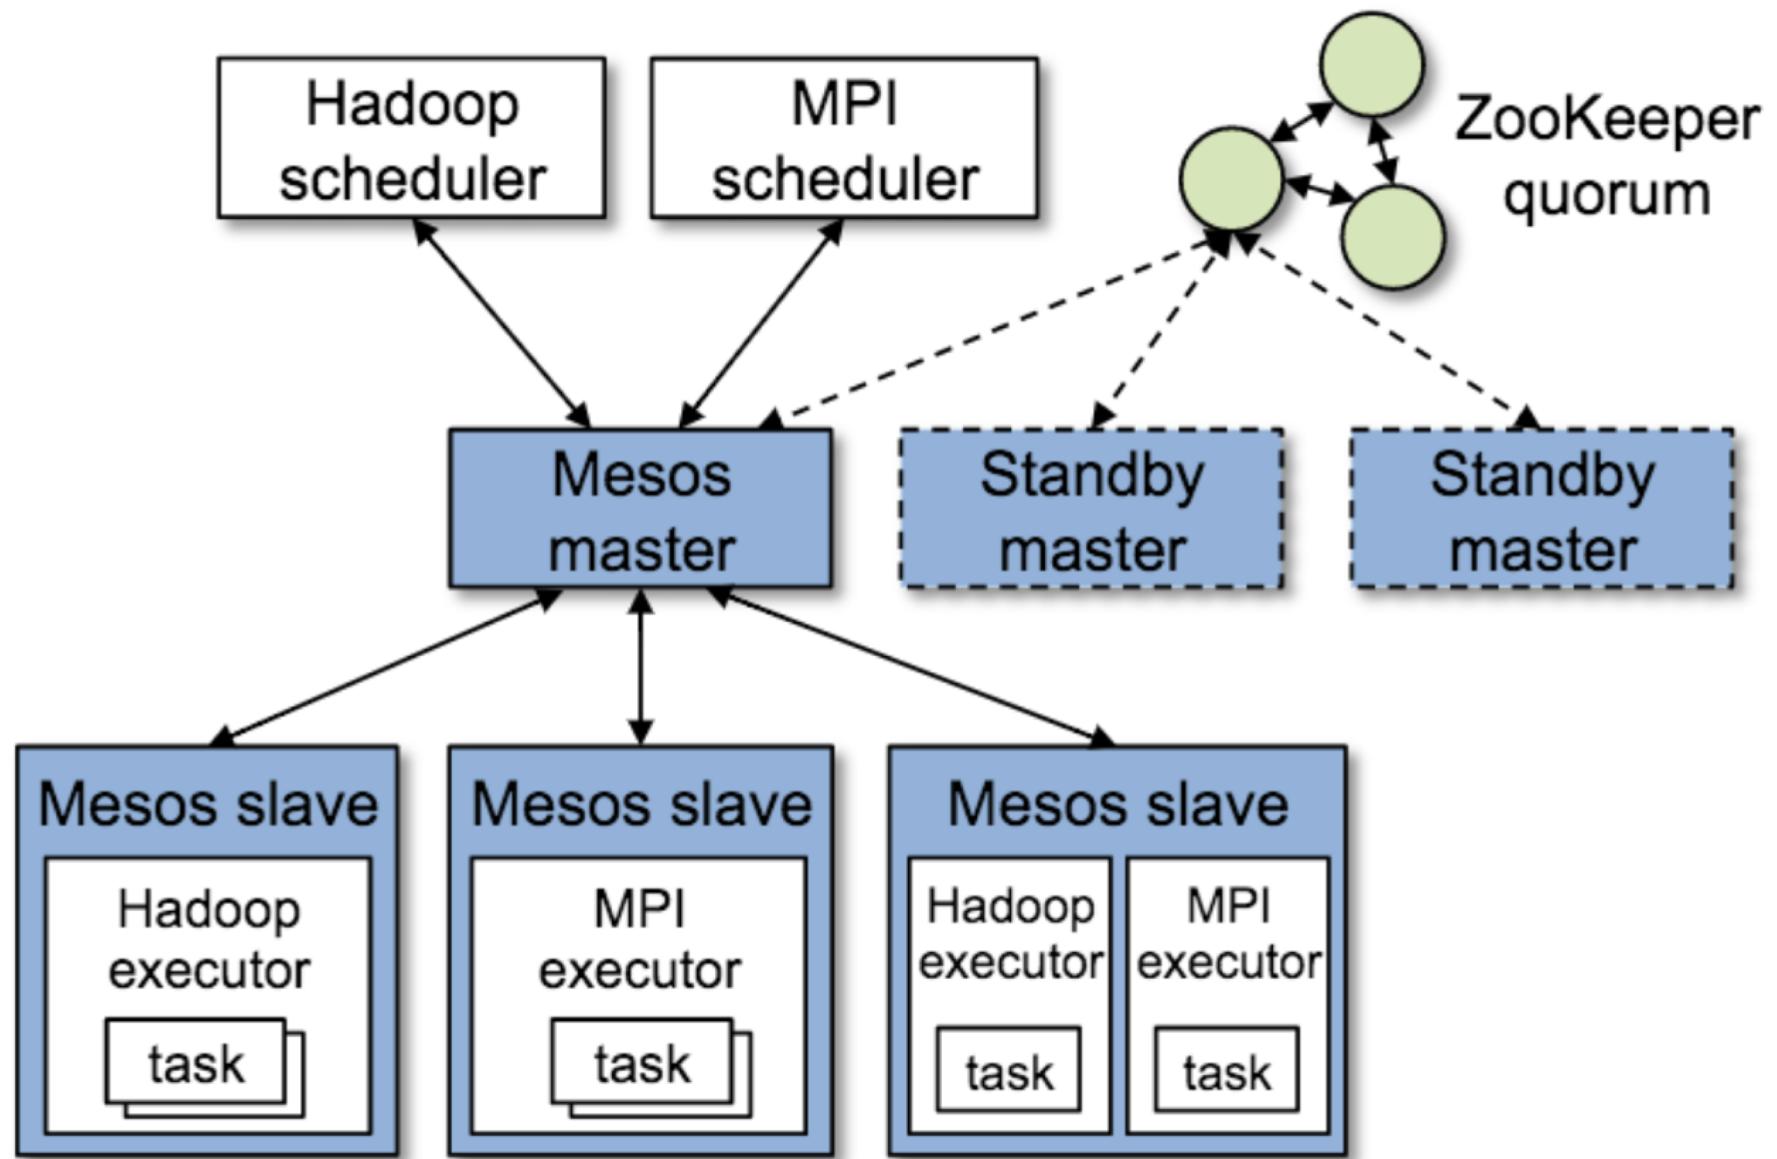
- **YARN is the system that runs your code on multiple nodes**
- Hadoop 2.0 replacement for the cluster manager
 - Basically a model to distribute and manage workloads
 - Not just MapReduce but supports other workloads



YARN architecture



Apache Mesos



Amazon Web Services

Compute

EC2

Virtual Servers in the Cloud

EC2 Container Service

Run and Manage Docker Containers

Elastic Beanstalk

Run and Manage Web Apps

Lambda

Run Code in Response to Events

Storage & Content Delivery

S3

Scalable Storage in the Cloud

CloudFront

Global Content Delivery Network

Elastic File System PREVIEW

Fully Managed File System for EC2

Glacier

Archive Storage in the Cloud

Import/Export Snowball

Large Scale Data Transport

Storage Gateway

Integrates On-Premises IT Environments with Cloud Storage

Database

RDS

Managed Relational Database Service

DynamoDB

Predictable and Scalable NoSQL Data Store

ElastiCache

In-Memory Cache

Redshift

Managed Petabyte-Scale Data Warehouse Service

Developer Tools

CodeCommit

Store Code in Private Git Repositories

CodeDeploy

Automate Code Deployments

CodePipeline

Release Software using Continuous Delivery

Management Tools

CloudWatch

Monitor Resources and Applications

CloudFormation

Create and Manage Resources with Templates

CloudTrail

Track User Activity and API Usage

Config

Track Resource Inventory and Changes

OpsWorks

Automate Operations with Chef

Service Catalog

Create and Use Standardized Products

Trusted Advisor

Optimize Performance and Security

Security & Identity

Identity & Access Management

Manage User Access and Encryption Keys

Directory Service

Host and Manage Active Directory

Inspector PREVIEW

Analyze Application Security

WAF

Filter Malicious Web Traffic

Internet of Things

AWS IoT BETA

Connect Devices to the cloud

Mobile Services

Mobile Hub BETA

Build, Test, and Monitor Mobile apps

Cognito

User Identity and App Data Synchronization

Device Farm

Test Android, Fire OS, and iOS apps on real devices in the Cloud

Mobile Analytics

Collect, View and Export App Analytics

SNS

Push Notification Service

Application Services

API Gateway

Build, Deploy and Manage APIs

AppStream

Low Latency Application Streaming

CloudSearch

Managed Search Service

Elastic Transcoder

Easy-to-use Scalable Media Transcoding

SES

Email Sending Service

SQS

Message Queue Service

SWF

Workflow Service for Coordinating Application Components

Enterprise Applications



EC2 / AWS main functions

- EC2 (Elastic Compute Cloud)
 - Instances
 - Servers of various sizes
 - AMIs (Amazon Machine Images)
 - Server images
 - Elastic Block Storage (EBS)
 - Virtualized Hard drives
 - VPC (Virtual Private Cloud)
 - Secure network space
- S3 (Simple Storage Solution)
 - “Buckets” of data
 - Longer term storage of data



Flintrock



license [Apache 2.0](#) build [passing](#) chat [on gitter](#)

Watch [@nchammas's talk](#) on Flintrock at Spark Summit East 2016: [talk](#) / [slides](#)

Flintrock is a command-line tool for launching [Apache Spark](#) clusters.



© Paul Fremantle 2015. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License
See <http://creativecommons.org/licenses/by-nc-sa/4.0/>

Installing standalone version

Standalone version (Python not required!)

If you don't have a recent enough version of Python, or if you don't have Python installed at all, you can still use Flintrock. We publish standalone packages of Flintrock on GitHub with our [releases](#).

Find the standalone package for your OS under our [latest release](#), unzip it to a location of your choice, and run the `flintrock` executable inside.

For example:

```
flintrock_version="0.11.0"

curl --location --remote-name "https://github.com/nchammas/flintrock/releases/download/v$flintrock_version/F
unzip -q -d flintrock "Flintrock-$flintrock_version-standalone-OSX-x86_64.zip"
cd flintrock/

# You're good to go!
./flintrock --help
```

You'll probably want to add the location of the Flintrock executable to your `PATH` so that you can invoke it from any directory.



Flintrock Launching a cluster

```
flintrock launch test-cluster \  
  --num-slaves 1 \  
  --spark-version 2.2.0 \  
  --ec2-key-name key_name \  
  --ec2-identity-file /path/to/key.pem \  
  --ec2-ami ami-a4c7edb2 \  
  --ec2-user ec2-user
```



Flintrock configure

```
▶ config.yaml x
1 services:
2   spark:
3     version: 2.2.0
4   hdfs:
5     version: 2.7.3
6
7 provider: ec2
8
9 providers:
10  ec2:
11    key-name: bigkp #for west-2
12    identity-file: /home/big/keys/bigkp.pem
13    instance-type: t2.micro
14    region: eu-west-2
15    ami: ami-01419b804382064e4 # Amazon Linux, eu-west-2
16    user: ec2-user
17    tenancy: default # default | dedicated
18    ebs-optimized: no # yes | no
19    instance-initiated-shutdown-behavior: terminate # terminate | stop
20
21 launch:
22   num-slaves: 2
23   install-hdfs: False
24
```



Other things you can do

```
flintrock destroy test-cluster
```

```
flintrock login test-cluster
```

```
flintrock describe test-cluster
```

```
flintrock add-slaves test-cluster  
    --num-slaves 2
```

```
flintrock remove-slaves test-cluster  
    --num-slaves 1
```

```
flintrock run-command test-cluster  
    'sudo yum install -y package'
```

```
flintrock copy-file test-cluster  
    /local/path /remote/path
```



Questions?



© Paul Fremantle 2015. This work is licensed under a Creative Commons
Attribution-NonCommercial-ShareAlike 4.0 International License
See <http://creativecommons.org/licenses/by-nc-sa/4.0/>