Big Data Engineering

Conclusions and Recap

Julie Weeds

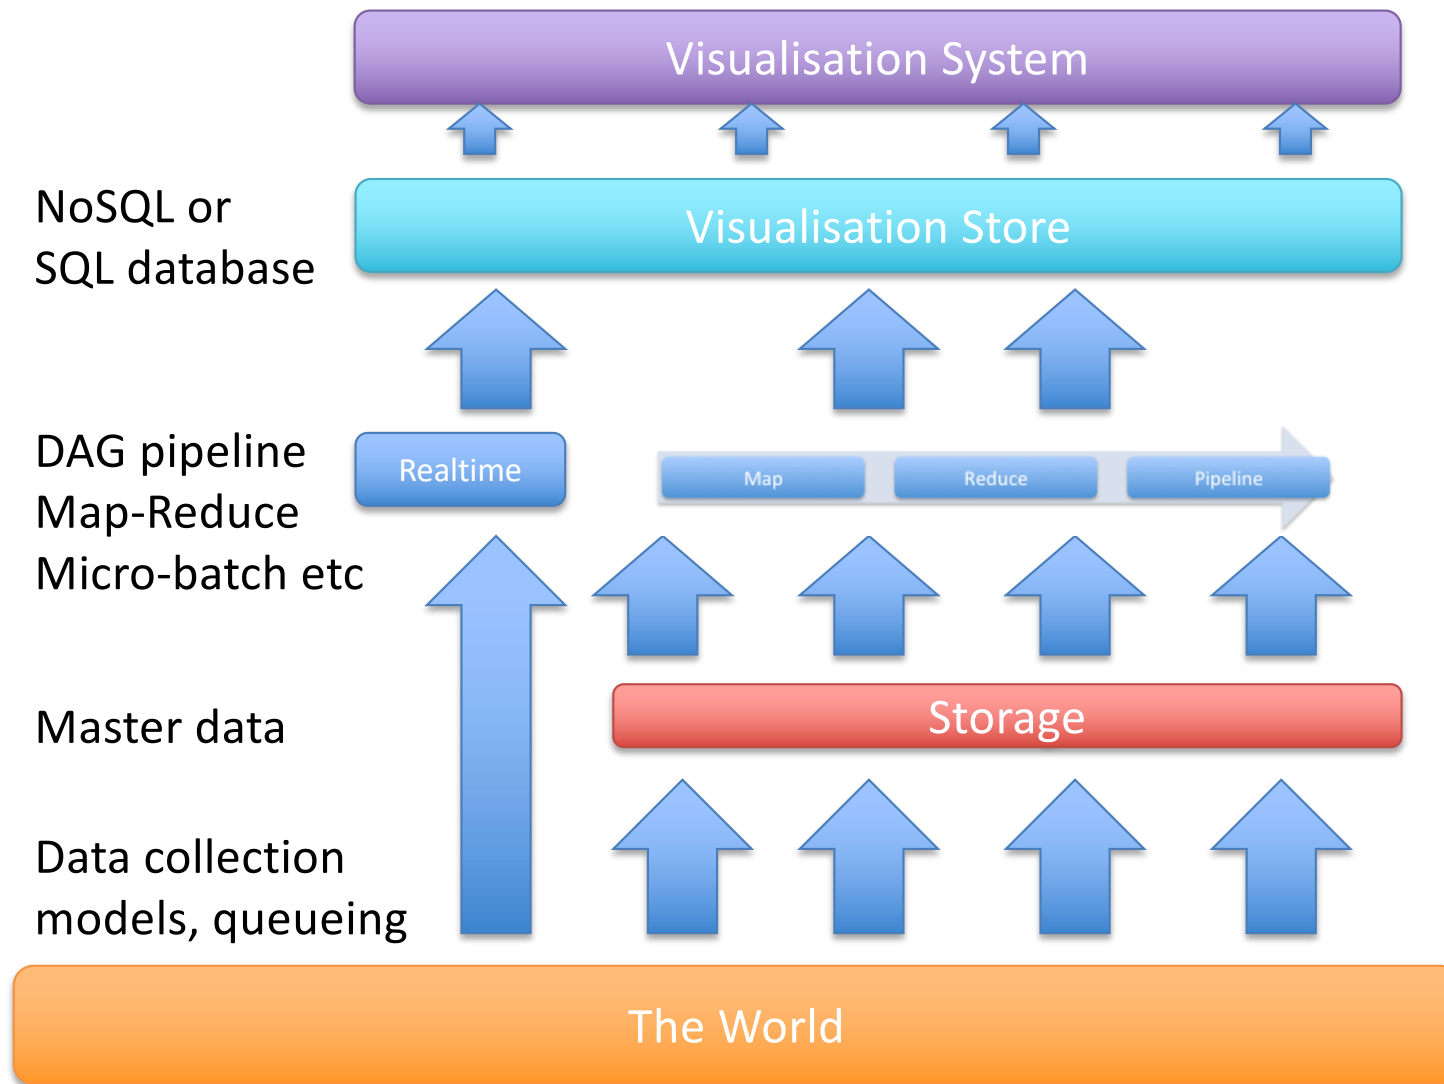March 2020

Contents

- Understanding the bigger picture

- What are the different components

- Message queueing and collection systems

- Map-Reduce and DAG systems

- Realtime Systems

- Fast databases for speed

- Visualisation and Dashboards

The big picture

Visualisation System

NoSQL or
SQL database

Visualisation Store

DAG pipeline
Map-Reduce
Micro-batch etc

Realtime

Map | Reduce | Pipeline

Master data

Storage

Data collection
models, queueing

The World

The big picture

- You have *immutable* master data

- You create a set of processes to:
  - Collect that data

  - Store master data

  - Process data

  - Visualise and present

- Some of those processes act on batch and others on real-time data

How to choose the components?

- Two main approaches:
  - Best of breed
    - Choose the best available component in each space
  - Stack
    - Choose a curated stack that a team or organization is providing/selling/supporting

Approach

- Minimise the pain
  - Choose what you need when you need it
  - Don't over engineer

How do I ingest data?

- File transfer

- Live stream
  - Sockets
  - Syslog
  - Messaging system

- From existing databases

How do I store data?

- HDFS

- zFS / GlusterFS / NFS …

- NoSQL database
    - Mongo / HBase / Cassandra / Apache Parquet

- CSV

How do I process data?

- Simple Map Reduce

- Hive / Pig

- DAG

- Pipeline

- etc

How do I visualise data

- From a SQL database?

- From a NoSQL database?

- Generate charts in Python Spark?

- Etc?

Collection / Queuing systems

- Two ways of making the choice
  - The protocol
  - The middleware
- Protocols
  - ZeroMQ, MQTT, AMQP, STOMP, Kafka Protocol, Rendevouz, etc
- Middleware
  - Kafka, Apollo, Mosquitto, QPid, WSO2, etc

Processing approaches

- Covered in detail already

- Hadoop

- Spark

- etc

Cluster Management

- Spark

- YARN

- Mesos

- Kubernetes

- etc

Visualisation

Visualisation approaches

- Full products
  - Tableau, Qlik, SAS, GoodData

- Web-based systems
  - Tableau Public, Datawrapper, Raw, Plotly

- Developer oriented
  - D3.js, dygraphs, Python charting, Leaflet, Fusion Charts, Google Charts, etc

Fortune top 10 big data companies

fortune.com/2014/06/13/these-big-data-companies-are-ones-to-watch/

- MapR – Apache Hadoop
- MemSQL
- Databricks – Apache Spark
- Platfora – Apache Hadoop
- Splunk
- Teradata – Apache Hadoop
- Palantir – Hadoop, Cassandra, Lucene
- Premise
- Datameer – Apache Hadoop
- Cloudera – Apache Hadoop
- Hortonworks – Apache Hadoop
- MongoDB – MongoDB
- Trifacta – Apache Hadoop

Trustradius big data companies to watch 2018

- Business Analytics
  - Alteryx
  - Arcadia Data
  - ClearStory Data
  - Cooladata
- Data management and integration
  - Actian
  - Alation
  - Attunity
  - Iguazio

- Big data platforms
  - **Hortonworks**
  - Micro Focus
  - **Teradata**
- Machine learning
  - DataRobot
  - H20.ai
  - **Splunk**
  - The Data Incubator
  - Domino

The real answer

You are on the cutting edge
- Expect to have some pain

Questions?